

RAID

Dr Peadar Grant

February 8, 2024

1 RAID

Redundant Array of Inexpensive Disks (RAID) is where a number of identical standard drives are used in multiple. Depending on the configuration used, this can increase capacity, performance and reliability.

RAID levels are standard patterns in which RAID is implemented, .

RAID set is a number of physical disks that has been combined using a RAID level.




Figure 1: RAID concept

1.1 Status

Healthy: where all of its member disks are working.

Degraded: where one or more of its member disks have failed while the others continue to work. I/O operations can continue.

Failed: where one or more of its member disks have failed such that I/O operations cannot continue.

2 RAID Metrics

2.1 Fault tolerance

Each RAID level can tolerate a number of disks failing.

2.2 Storage efficiency

The storage efficiency, E , of a RAID set is:

$$E = \frac{\text{useable capacity}}{\text{sum of RAID set disks capacity}} \quad (1)$$

For any given RAID set, it depends on the RAID level, L , in use and the number of disks, n .

2.3 Write penalty

Most RAID levels have a write penalty, which is the number of write operations that occur compared to when a single disk is used. A single disk would have a write penalty of 1.

3 Basic RAID levels

The two basic RAID levels involve **striping** and **mirroring**. Assume that:

- data has been split up into equally sized chunks labelled A1 , A2 etc.
- there are n disks in the set.

3.1 RAID-0 (Stripe)

RAID-0 splits, or stripes, data across a minimum of 2 disks, Figure 2.

3.1.1 RAID-0 fault tolerance

RAID-0 has **no fault tolerance**.

- If any one disk fails no I/O can be performed.
- If the failure is permanent, data is permanently lost.

3.1.2 RAID-0 write penalty

RAID-0 does not incur additional writes, so its write penalty is 1.




Figure 2: RAID-0 with 2 disks

3.1.3 RAID-0 storage efficiency

In RAID-0, data is striped across all disks in the set:

- The space of all the n disks in the set is usable.

$$\text{RAID-0 efficiency} = \frac{n}{n} = 1 \quad (2)$$

3.2 RAID-1 (Mirror)

RAID-1 mirrors, or copies, data across n disks:

- Usually 2 disks, can have more.
- Read performance theoretically n times single disk.
- Write performance is identical to a single disk in theory.

3.2.1 RAID-1 fault tolerance

RAID-1 maintains full data integrity with only a single disk remaining functional:

- RAID-1 disk is identical to a single non-RAID disk.




Figure 3: RAID-1

3.2.2 RAID-1 write penalty

- RAID-1 incurs a write for each disk in the set, so its write penalty is n .
- Most RAID-1 sets have two disks, so this is often 2.

3.2.3 RAID-1 storage efficiency

In RAID-1, data is mirrored onto all disks in the set. So only 1 disk's worth of capacity in the set of n is useable, giving:

$$\text{Efficiency}(L=1, n) = \frac{1}{n} \quad (3)$$

Example 1: RAID-1 storage efficiency

Calculate the storage efficiency of a 2-disk RAID-1 set.

$$\text{Efficiency}(L=1, 2) = \frac{1}{2} \quad (4)$$

$$= 0.5 \quad (5)$$

4 Parity-based RAID

Parity relies on the XOR operation:

A	B	$C = A \oplus B$
0	0	0
0	1	1
1	0	1
1	1	0

Table 1: XOR operation

The parity blocks are computed bitwise from the data blocks using XOR.

4.1 RAID-4

RAID-4 stripes data blocks across $n-1$ disks:

- remaining disk used for parity
- Minimum 3 disks.

4.1.1 Parity computation

Bit-by-bit the parity computation for a single stripe, here A is:

$$A_p = A_1 \oplus A_2 \oplus A_3 \quad (6)$$

Same for stripes B, C, D .




Figure 4: RAID-4

4.1.2 RAID-4 storage efficiency

For RAID-4, the parity data takes up the capacity of the one parity disk:

- The parity data does not contribute to the storage available for use.
- Of the n disks in the set, only the capacity of $n-1$ disks is available for use:

$$\text{Efficiency}(L=4, n) = \frac{(n-1)}{n} = 1 - \frac{1}{n} \quad (7)$$

Example 2: RAID-4 storage efficiency

Calculate the storage efficiency of a 3-disk RAID-4 set.

$$\text{Efficiency}(L=4, 3) = \frac{(3 - 1)}{3} \quad (8)$$

$$= \frac{2}{3} \quad (9)$$

$$= 67 \% \quad (10)$$

4.2 RAID-5

RAID-5 is conceptually similar to RAID-4:

- Parity is computed for each stripe in the same way, but is stored differently.
- Instead of a single parity disk, the parity blocks are distributed amongst all disks.




Figure 5: RAID-5

4.2.1 **RAID-5 fault tolerance**

- RAID-5 is tolerant of failure of one disk. The array can continue to operate in degraded mode with no data loss and I/O operations available.
- When the failed disk is replaced, a rebuild occurs onto the new disk by XORing the corresponding blocks on the remaining disks.
- The array is vulnerable while degraded. If a second disk failure occurs, data will be lost.

4.2.2 RAID-5 storage efficiency

For RAID-5, the parity data takes up the capacity of one disk (although distributed across them all):

- The parity data does not contribute to the storage available for use.
- This means that of the n disks in the set, only the capacity equivalent to $n-1$ disks is available for use:

$$\text{Efficiency}(L=5, n) = \frac{(n - 1)}{n} = 1 - \frac{1}{n} \quad (11)$$

Example 3: RAID-5 storage efficiency

Calculate the storage efficiency of a 4-disk set in RAID-5.

$$\text{Efficiency}(L=5, 4) = \frac{(4 - 1)}{4} \quad (12)$$

$$= \frac{3}{4} \quad (13)$$

$$= 75\% \quad (14)$$

4.3 RAID-6

RAID-6 is similar to RAID-5 but adds a second parity block in each stripe.

4.3.1 RAID-6 fault tolerance

The second parity block in each stripe means that RAID-6 can sustain the loss of up to two disks in a set while still remaining operational:

- Rebuild proceeds in the same way as RAID-5 except there are two parity blocks.




Figure 6: RAID-6

4.3.2 RAID-6 storage efficiency

In RAID-6 we lose the equivalent of two disks' space to store the parity information.

$$\text{Efficiency}(L=6, n) = \frac{n-2}{n} = 1 - \frac{2}{n} \quad (15)$$

Example 4: RAID-6 storage efficiency

Calculate the storage efficiency of a 4 disk RAID-6 set.

$$\text{Efficiency}(L=6, 2) = \frac{4-2}{4} \quad (16)$$

$$= 0.5 \quad (17)$$

5 Nested RAID

RAID 0+1

RAID 1

RAID 0

RAID 0




Figure 7: RAID-01

RAID 1+0

RAID 0

RAID 1

RAID 1




Figure 8: RAID-10

6 Implementation

RAID can be implemented in either software or hardware:

Hardware-based RAID is where the Host Bus Adapter (HBA) implements the RAID functionality. The host and its OS sees the single block device resulting from the RAID implementation.

Software-based RAID is where the RAID functionality is implemented within the host, usually as part of the operating system. The HBA just passes through each physical disk as a block device.

Hybrid RAID is sometimes used for nested RAID configurations. Disks are combined into RAID sets by one or more HBAs and the exported block devices are then grouped into RAID set(s) in software.

Hardware-assisted RAID is a mix of hardware and software RAID. The XOR computations are accelerated by a chip in the HBA while the rest of the RAID implementation is done in software. Seen in some motherboard-based RAID controllers and often best avoided!

Despite a lot of ill-informed comment online and elsewhere, it is almost impossible to definitively say whether software or hardware-based RAID is preferable. The choice depends on a large number of factors: devices, HBA capabilities, HBA throughput, OS, filesystem(s) to be used, portability requirements.

7 Monitoring

Problem: RAID set becomes degraded due to disk failure. Essential that the failed disk is replaced before a further failure occurs.

Solution: The RAID set must have appropriate monitoring installed.

- Hardware RAID controllers often have a separate management capabilities, appearing as a separate device on the host system. This often needs proprietary software to work.
- Software RAID setups more visible by default to the OS and will appear in system logs.
- Either way, should be proactively sending notifications (e.g. e-mail, helpdesk API).

8 Hot spares

If hardware / software capabilities and available drive bays allow, it is possible to automatically replace a failed disk with a hot spare.

1. The RAID controller / software detects the failed disk, sends notification.
2. Controller removes it from the array (which is now degraded). The disk remains physically in its bay and plugged into the same port.
3. The controller/software then adds the hot spare to the array. Re-build commences on to the hot spare.
4. Array is then in a healthy state once rebuild has completed.
5. Failed disk can be replaced on schedule on next physical visit. Re-built onto replaced disk, hot spare removed from array.