

Filesystems

Dr Peadar Grant

February 14, 2024

Contents

1	Disk-based filesystems	S.2
2	Journalling	S.5
3	Single vs multi-root	S.8
4	Virtual filesystems	S.19

1 **Disk-based filesystems**

Disk-based file systems interposes between the application layer (including the user-space OS tools like File Manager) and the block devices. A file system is created on a block device (= a block device is formatted with a file system) that is then mounted by the host's operating system.

1.1 Common filesystems

As of the present time the most common disk-based filesystems are:

Microsoft Windows: normally NT File System (NTFS) for fixed disks.

Apple Mac OS X: APFS (solid-state disks only), HFS+ (formerly all fixed media, remains on magnetic disks only)

Linux-based systems commonly use ext4, but a large range of other systems commonly appear: ext3, xfs, reiserfs, ZFS*

UNIX-based systems tend to have an OS-specific filesystem w/ similar behaviour to linux, often ZFS

Wikipedia has collated a comparison of file systems.

1.2 Roles of filesystems

Space management using allocation table. Dealing with *fragmentation*.

Naming possibly case (sensitive | insensitive | preserving)

Hierarchy using directories.

- Historically and some specialist systems: **flat** filesystem

Metadata such as file creation / modification time

Permissions management via metadata to allow (read | write | execute) to specific users.

- Common permissions are owner / group / world (others) in UNIX.
- The host OS, not the filesystem, actually enforces these restrictions.

Maintaining integrity using redundancy, checksums.

Durable storage by journalling

2 Journaling

Durability is the key expectation:

- data once written to filesystem isn't lost (within reasonable assumptions).

Journaling is where the intent to change data is recorded to the journal before the actual data itself is modified:.

1. Data is written to the on-disk journal
2. The actual on-disk file data is changed
3. The journal is deleted

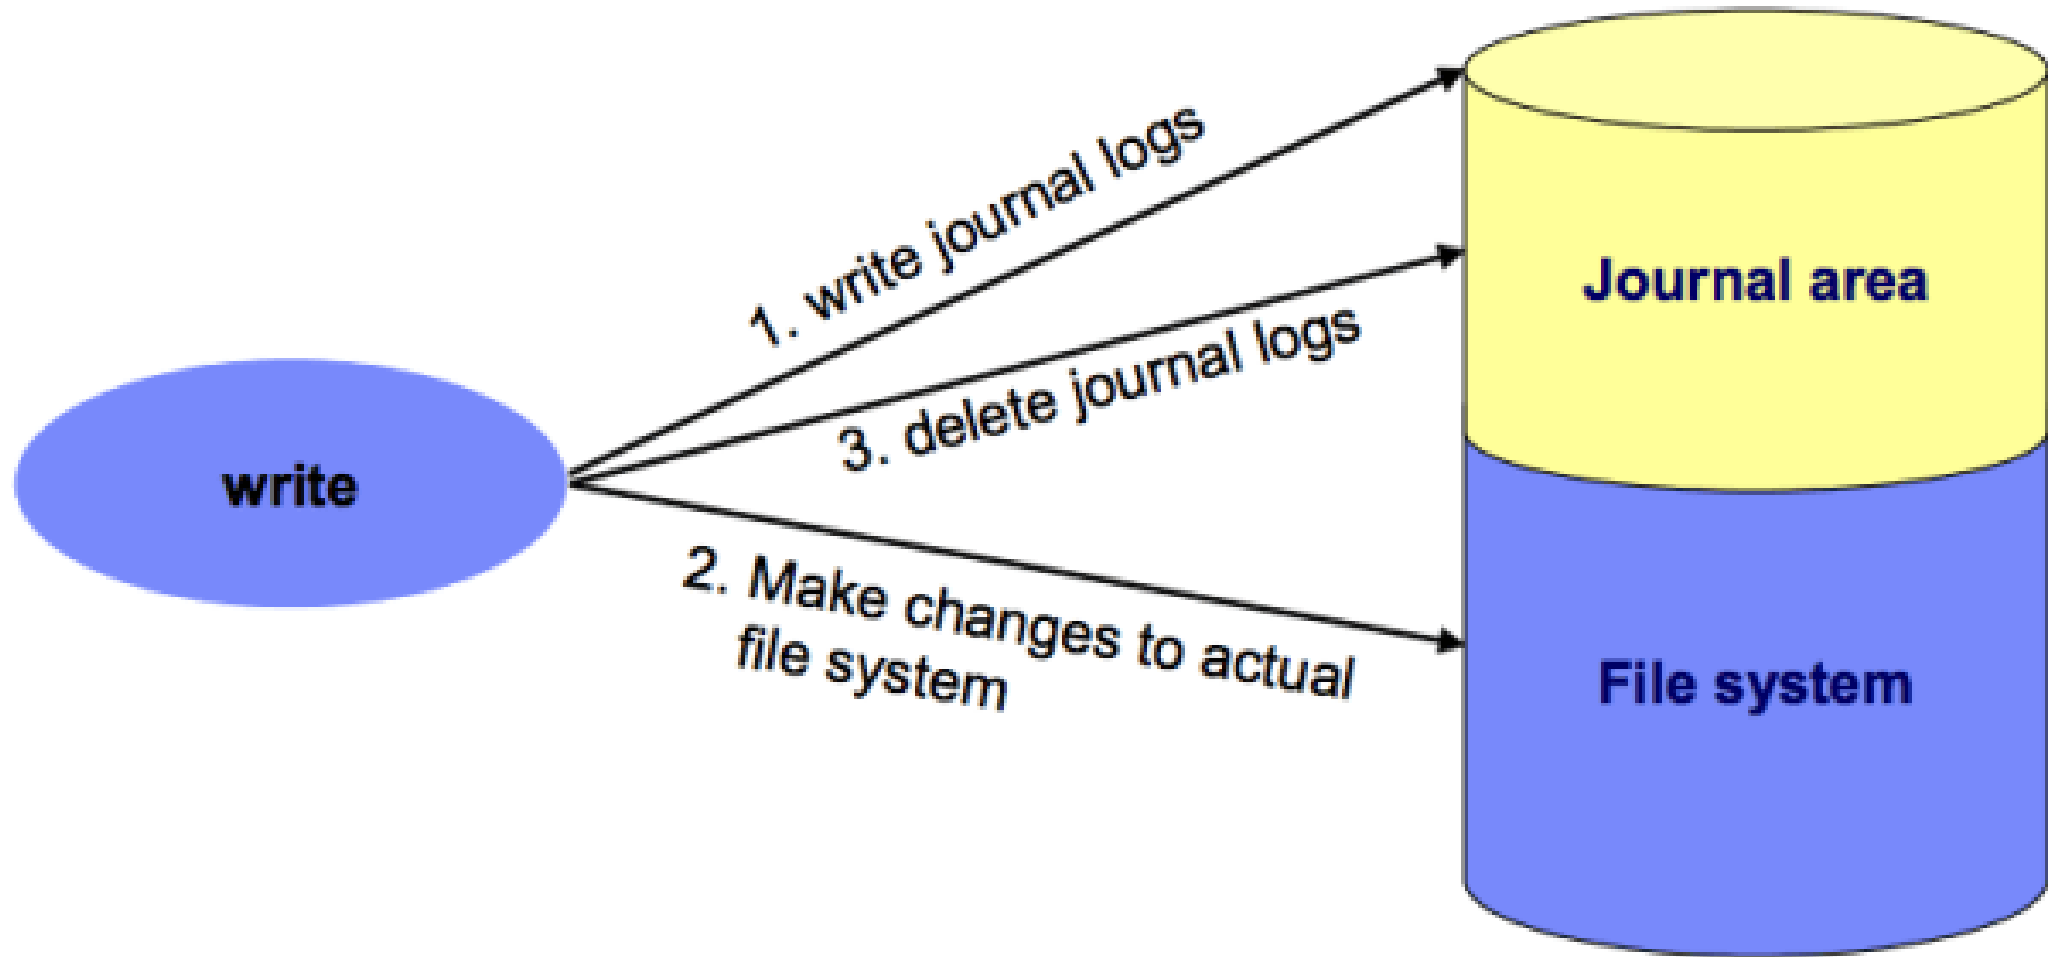


Figure 1: Journalling filesystem

2.1 Alternatives to Journalling

Copy-on-write: where a modification is first written and then the original deleted

Log-structured file system: the journal itself is the filesystem

3 Single vs multi-root

Although more an OS-characteristic, there is an important difference between single and multi-root operating systems that becomes relevant for discussion of filesystems:

Single root operating systems

- boot with **one** filesystem mounted as the **root** / directory.
- All others are attached at *mountpoints* within the root directory.

Multi-root operating system has multiple roots:

- often indicated by drive-letters e.g. C: , D:.
- other OSes (e.g. VMS, Novell) use volume labels (e.g. SYS)

In general, most UNIX-like OS are single root. Windows is multi-root.

3.1 Linux file system operations

Linux and UNIX tend to use similar commands to those shown here:

- Always check the precise requirement for your OS.

3.2 Creation

File system creation, also called Formatting, is done through the `mkfs` command:

```
# review help text and manpage
```

```
mkfs -h
```

```
man mkfs
```

```
# check where mkfs is located (usually /sbin)
```

```
which mkfs
```

```
# from that, find the available filesystems
```

```
ls -l /sbin | grep mkfs
```

```
# creating a filesystem (aka formatting)
```

```
# e.g. make an ext4 on /dev/sdb1
```

```
mkfs -t ext4 /dev/sdb1
```

3.3 Mounting

Mounting is where a particular a block device formatted with a particular filesystem is attached into the file system hierarchy:

- A folder, called the “mount point” is created on an existing (usually the root) filesystem.
- Mounting the additional drive then “covers up” that folder with root of the mounted drive.

Consider the example of mounting a block device `/dev/sdb1` at the mountpoint `/mnt/data`:

```
# review mount command
```

```
mount -h #inbuilt help
```

```
man mount #man page
```

```
# create the mountpoint (-p allows it to already exist)
```

```
mkdir -p /mnt/data
```

```
# issue the mount}
```

```
mount -t ext4 /dev/sdb1 /mnt/data
```

```
# viewing mounted systems
```

```
mount
```

3.4 Mounting on boot

The manual method has some shortcomings:

- Manual mounts / unmounts persist only as long as the system is running.

Instead we'd like to have a configuration file of mountable filesystems.

The `/etc/fstab` file provide this:

- lists filesystems that can be mounted / unmounted without specifying the block device to the `mount` command
- additionally can indicate these should be mounted automatically when the system starts up.

3.5 **/etc/fstab** format

The `/etc/fstab` file format is whitespace separated of the following fields:

1. device-spec of the device to mount (device name, label, UUID, other means to identify)
2. mount-point where filesystem can be accessed once mounted
3. fs-type identifying the filesystem
4. options that may be needed: defaults per-filesystem, noauto to not mount on boot
5. dump - for use by the dump program for backups (backups to be discussed later)
6. pass - order for fsck to check errors at boot time: 0=don't, 1=during, 2=after

More advanced generic example, on Wikipedia:

# device-spec	mount-point	fs-type	options
LABEL=/	/	ext4	defaults
/dev/sda6	none	swap	defaults
none	/dev/pts	devpts	gid=5,mode=620
none	/proc	proc	defaults
none	/dev/shm	tmpfs	defaults
# Removable media			
/dev/cdrom	/mnt/cdrom	udf,iso9660	noauto,owner,ro
# NTFS Windows 7 partition			
/dev/sda1	/mnt/Windows	ntfs-3g	quiet,defaults,locale=en_US.utf8,umask=0,r
# Partition shared by Windows and Linux			
/dev/sda7	/mnt/shared	vfat	umask=000


```
# Mounting tmpfs
tmpfs          /mnt/tmpfschk  tmpfs          size=100m

# Mounting cifs
//cifs_server_name/ashare  /store/pingu  cifs          credentials=/root/smbpass.txt

# Mounting NFS
nfs_server_name:/store    /store        nfs           rw
```

Simpler sample of /etc/fstab on a Raspberry-Pi:

```
proc          /proc          proc          defaults      0            0
PARTUUID=af832fc2-01  /boot          vfat          defaults      0            2
PARTUUID=af832fc2-02  /              ext4          defaults,noatime 0            1
/dev/sdb2            /data         ext4          defaults,noatime 0            0
# a swapfile is not a swap partition, no line here
```

```
# use dphys-swapfile swap[on|off] for that
```

Can often also identify block devices by other means than their /dev node, for example partition UUID. Avoids ambiguities if drives are swapped around. Check help for mount command to be sure.

3.5.1 Automounters

Other means to identify and mount filesystems exist on Linux (particularly) nowadays: systemd unit files, pmount, automount.

3.6 Windows

3.6.1 Creation

By demonstration through Computer Management.

3.6.2 Mounting

Generally a drive letter will be assigned.

Windows can also do a UNIX-style mount to a folder within another filesystem. Not commonly used though.

4 Virtual filesystems

A filesystem as such is implemented within the operating system normally in the form of code. This may in fact allow filesystems to be created that do not directly reside on any block device.

Network filesystem where the host connects to a filesystem located on another computer using file-sharing protocols like SMB, NFS.

Object storage filesystems where the host connects to a storage service that does not itself implement the normal hierarchical filesystem, but can appear to do so:

- Software to map WebDAV and systems like S3

Device filesystem usually mounted at /dev used in UNIX-like OSes to provide access and permission control to peripheral devices (e.g. HID, serial, LPT, block devices).

Special filesystems seen in UNIX:

- sysfs and configfs kernel

- /proc within info on processes